

# Yuqing Zhou

✉ yzhou31@gmu.edu | 🌐 yuqing-zhou.github.io

## Research Interests

---

Self-evolving AI, LLM Reasoning, LLM Post-training, Trustworthy AI

## Education

---

**George Mason University**

*Ph.D. in Computer Science*

**Aug. 2023 – Present**

*Fairfax, VA*

**University of Michigan**

*M.S. in Electrical and Computer Engineering, Computer Vision*

**Aug. 2019 – Apr. 2021**

*Ann Arbor, MI*

**Southeast University**

*B.Eng. in Electronic Science and Technology*

**Aug. 2015 – Jun. 2019**

*Nanjing, China*

## Research & Industry Experience

---

**Applied Scientist Intern, Amazon Web Services (AWS)**

*Efficient LLM Reasoning*

**May 2026 – Aug. 2026**

*Seattle, WA*

- Developed a step-aware credit assignment method for RL-based post-training that redistributes trajectory-level learning signals according to the contribution of individual reasoning steps, preserving useful computation while suppressing redundant and dead-end reasoning.
- On Qwen2.5-Math-7B, outperformed trajectory-level length regularization, trace-compression, and fine-grained credit-assignment baselines, reducing reasoning tokens by **24% on average** while maintaining or improving accuracy.

**Applied Scientist Intern, Amazon Web Services (AWS)**

*Agentic Conversational AI*

**May 2025 – Aug. 2025**

*Seattle, WA*

- Developed workflow-based control mechanisms for tool-using conversational agents, representing procedures as structured memory with prerequisite-aware action specifications.
- Built an offline-to-online learning pipeline that induces workflows from historical agent trajectories and uses runtime progress alignment to guide agent execution.
- Evaluated the approach on interactive agent benchmarks including ToolSandbox and  $\tau^2$ -bench, consistently improving agent performance and execution reliability over baseline agents.

**Software Engineer, Shanghai Huawei Technologies Co., Ltd,**

*Shanghai, China*

**Aug. 2021 – May 2023**

## Publications

---

- [1] Bowen Wei, Nan Wang, **Yuqing Zhou**, Jinhao Pan, and Ziwei Zhu. “Confidence-Orchestrated Self-Evolution against Uncertain LLM Feedback”. In: *EMNLP* (2026).
- [2] **Yuqing Zhou** and Ziwei Zhu. “Fighting Spurious Correlations in Text Classification via a Causal Learning Perspective”. In: *NAACL* (2025).
- [3] **Yuqing Zhou**, Ruixiang Tang, Ziyu Yao, and Ziwei Zhu. “Navigating the Shortcut Maze: A Comprehensive Analysis of Shortcut Learning in Text Classification by Language Models”. In: *Findings of EMNLP* (2024).
- [4] **Yuqing Zhou**, Tianshu Feng, Mingrui Liu, and Ziwei Zhu. “A Generalized Propensity Learning Framework for Unbiased Post-Click Conversion Rate Estimation”. In: *CIKM* (2023).

## Selected Research Projects

---

### Causally Calibrated Robust Classifier (CCR)

Mar. 2024 – Oct. 2024

- Developed a group-label-free framework for improving worst-group robustness under distribution shift by reducing model reliance on spurious correlations.
- Designed a causal feature selection mechanism together with inverse propensity weighting to identify and downweight unstable predictive correlations without requiring group annotations.
- Achieved **70.7% worst-group accuracy** on CivilComments, outperforming competing robust-learning methods that require explicit group labels.

### Shortcut Learning Analysis in Language Models

Sep. 2023 – Jun. 2024

- Developed a systematic taxonomy of shortcut learning in language models, covering occurrence-, style-, and concept-based shortcuts.
- Constructed controlled benchmarks for these failure modes, including LLM-generated datasets with synthetic shortcut correlations.
- Evaluated BERT and Llama-family models, including Llama2-7B/13B and Llama3-8B, together with state-of-the-art robust learning methods, showing that increased model scale alone does not consistently improve shortcut robustness.

### Generalized Propensity Learning for Post-Click Conversion Rate Prediction

May 2023 – Jun. 2023

- Developed a generalized propensity-learning (GPL) framework to minimize bias and variance in post-click conversion prediction (CVR) for recommender systems, enhancing the performance and robustness of existing methods like IPS and DR-based estimators.
- Improved CVR prediction accuracy by **7% in DCG@2** and **6% in Recall@2** on the Yahoo dataset.

## Technical Skills

---

**Programming:** Python, C/C++, Julia, MATLAB

**ML Frameworks:** PyTorch, Hugging Face Transformers

**LLM Training & Inference:** verl, TRL, vLLM

## Teaching Experience

---

### Graduate Teaching Assistant, George Mason University

CS 657: Mining Massive Datasets

Fall 2026

CS 584: Theory and Applications of Data Mining

Fall 2026

CS 455: Computer Communications and Networking

Fall 2024, Spring 2025

## Academic Service

---

**Reviewer:** ICLR (2025–2027), KDD (2026–2027), PAKDD (2025), AISTATS (2025), IEEE TKDE, ACM TOIS, ACM TIST, ACM TORS